

Predicting Cannabis Consumption from Demographics and Personality

HarvardX PH125.9x Data Science Capstone

Charles Mégnin

10/8/2019

Contents

Executive Summary	1
+ Introduction	1
+ Goal of project	1
+ Dataset description	2
+ Key steps	2
Analysis	4
A: Data engineering	4
B: Data exploration	5
C: Modeling	12
Results	15
Conclusion	15

Executive Summary

+ Introduction

Drug use is a behavior that constitutes an important factor linked to poor health, including early mortality, and which presents significant adverse consequences for the social fabric, notably with respect to criminality and family cohesion. Early detection of an individual's predisposition to drug consumption offers healthcare professionals an opportunity to short-circuit the onset of addiction.

The present study is based on a dataset that includes demographic and psychological information related to the consumption of 18 legal and illegal drugs by 1885 participants. For the purpose of this study, we choose to focus the data analysis and modeling on the use of cannabis.

+ Goal of project

The goal of this project is to assess whether an individual's consumption of cannabis can be predicted from a combination of demographic and personality data.

To do so, we build and assess the effectiveness of six machine learning classifiers and confront the results obtained with the insights provided by data exploration.

+ Dataset description

The original dataset is found on the UCI machine learning repository. It is based the research paper by E. Fehrman, A. K. Muhammad, E. M. Mirkes, V. Egan and A. N. Gorban, “The Five Factor Model of personality and evaluation of drug consumption risk.,” arXiv, 2015. The data was collected from 1885 English-speaking participants over 18 years of age between March 2011 and March 2012.

In the original dataset, drug use is separated between ‘Never used’, ‘Used over a decade ago’, ‘Used over a decade ago’, ‘Used in last decade’, ‘Used in last year’, ‘Used in last month’, ‘Used in last week’ and ‘Used in last day’. For the purpose of this study, we separate the data in two groups: ‘Never Used’ (the original predictor) and ‘Used’ (the combination of the others, representing people having used cannabis at least once in their lifetime).

The original dataset includes answers to questions related to the use of alcohol, amphetamines, amyl nitrite, benzodiazepines, cannabis, chocolate, cocaine, caffeine, crack, ecstasy, heroin, ketamine, legal highs, LSD, methadone, magic mushrooms, nicotine and volatile substance abuse (VSA)) and one fictitious drug (Semeron) which was introduced to identify over-claimers. In the present study, we restrict our scope to the analysis of cannabis consumption.

The data consists of two groups of pre-normalized and centered predictors:

1. Five demographic predictors : Age, Gender, Level of education, Ethnicity, and Country of origin.
2. The results from seven scored tests administered to assess personality, specifically:
 - Neuroticism (a long-term tendency to experience negative emotions such as nervousness, tension, anxiety and depression);
 - Extraversion (manifested in outgoing, warm, active, assertive, talkative, cheerful, and in search of stimulation characteristics);
 - Openness to experience (a general appreciation for art, unusual ideas, and imaginative, creative, unconventional, and wide interests);
 - Agreeableness (a dimension of interpersonal relations, characterized by altruism, trust, modesty, kindness, compassion and cooperativeness);
 - Conscientiousness (a tendency to be organized and dependable, strong-willed, persistent, reliable, and efficient);
 - Impulsiveness;
 - Sensation-seeking.

The working dataset in this study therefore consists of one Class (Cannabis consumption labeled ‘Used’) and twelve predictors (five demographic and seven personality-related).

+ Key steps

We extract a training subset (80% of data) from the dataset for the purpose of training our model, and use the remaining 20% of the data as a test set for the purpose of evaluation. This being a classification problem, we use accuracy as the metric to assess the goodness of fit.

This report consists of two main sections:

- In the first part, after performing minor data engineering, we explore, bin, and analyze the dataset.
- In the second part, we move on to the modeling phase:
 - After applying a Recursive feature elimination algorithm to seek and discard predictors that do not contribute significantly to the outcome, we build models based on the following methods:
 - Generalized linear model (glm)
 - Generalized linear model with penalized maximum likelihood (GLMnet)

- Decision tree (rpart)
- Random forest (rf)
- Stochastic gradient boosting (gbm)
- Neural network (nnet)

We compare the modeling approaches, both in terms of accuracy and coherence with the data analysis.

Analysis

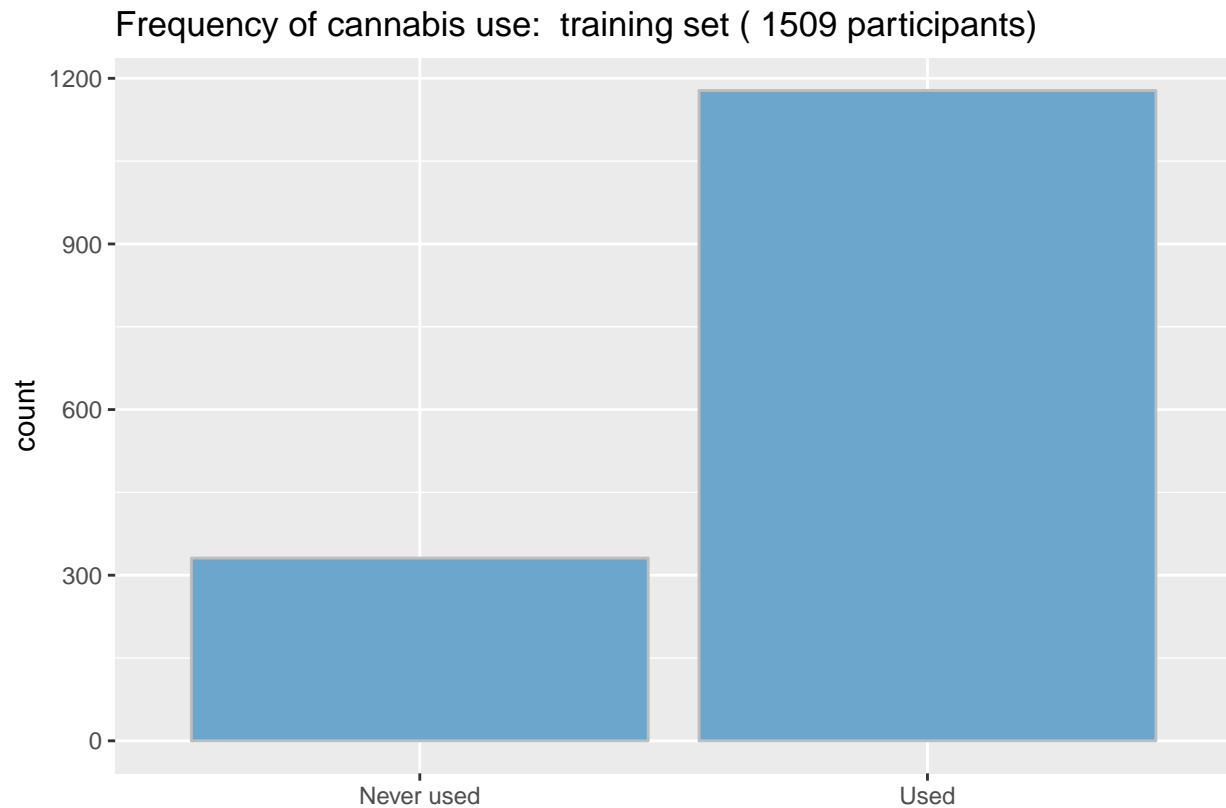
A: Data engineering

- All predictors were already normalized and centered in the original dataset.
- We construct the ‘Used’ class to separate ‘Never used’ participants (0) from the others (1).
- We then partition the data between training (80% / df.train) and test sets (20% / df.test) preserving the distribution of the Cannabis class.

B: Data exploration

- There are 0 NAs in the dataset.



Class distribution






The training set consists of 1178 participants having used cannabis at some point in their lives and 331 participants who haven't for a user-to-non-user ratio of 1: 3.6.

Contingency plots (prior to binning)

Cannabis use by demographic group







We note that the dataset of 1509 participants is dominated by young and educated white American and British participants of both sexes.


Binning

The small size of many demographic sub-groups add little valuable insight and will likely only serve to introduce variance in the analysis. At the risk of erasing behavioral differences among groups, the distribution of the dataset forces a more meaningful binning of the demographic information:


- 5 age groups: “18-24”, “25-34”, “35-44”, “45-54”, “55+”
- 5 groups for Education: “Left school as a teen”, “Some college”, “Professional certificate”, “University degree”, “Graduate degree”.
- 3 groups for Country: “USA”, “UK”, “Others”.
- 2 ethnic groups: “Whites”, “Non-whites”

Analysis of demographics

Use of cannabis in training set by:



Personality analysis




The density plots show some measure of difference between users and non-users particularly as it relates to either openness to experience, agreeableness, conscientiousness, impulsivity, and sensation-seeking. Some implications are rather entertaining, notably the notion that nice (ie: agreeable) people may be less likely to smoke cannabis, or conversely that exposure to pot might make people less nice.

On the other hand, the distributions for users and non-users as they relate to neuroticism or extraversion are similar, suggesting that these personality traits may not impact cannabis consumption. We will examine in the modeling section below whether that is indeed the case.

Feature correlation

We examine redundancies among the 12 predictors and with the Used class:




Besides the goodness of fit to the test data, the demographic and personality-related observations above will guide the assessment of the models we derive.

C: Modeling


We seek a model which improves on the ratio of users to the population (78.10%). This constitutes the baseline above which predictive modeling is interesting.

Recursive Feature elimination

For RFE as well as subsequent modeling, we use the k-fold cross validation method which involves splitting the dataset into k subsets. The algorithm holds aside one of the subsets while the model is trained on the others. This process is repeated a predetermined number of times and the overall accuracy estimate is provided.

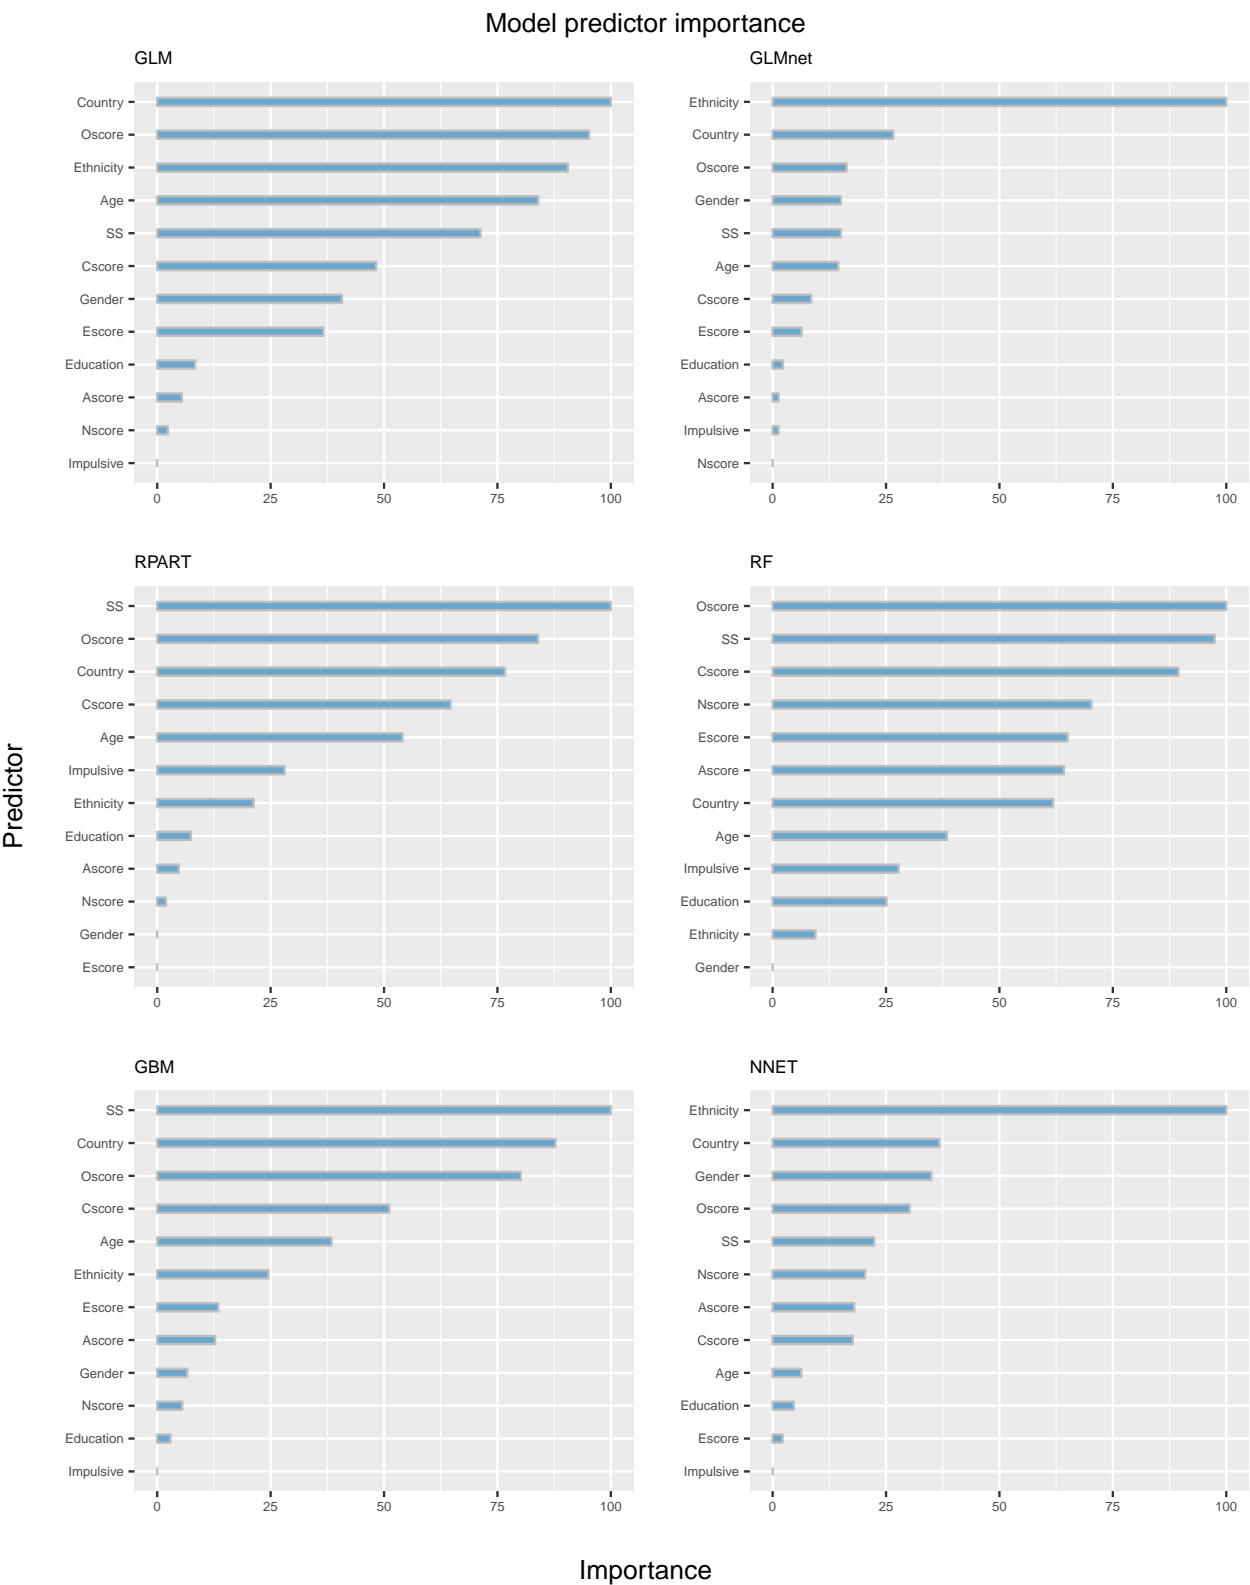


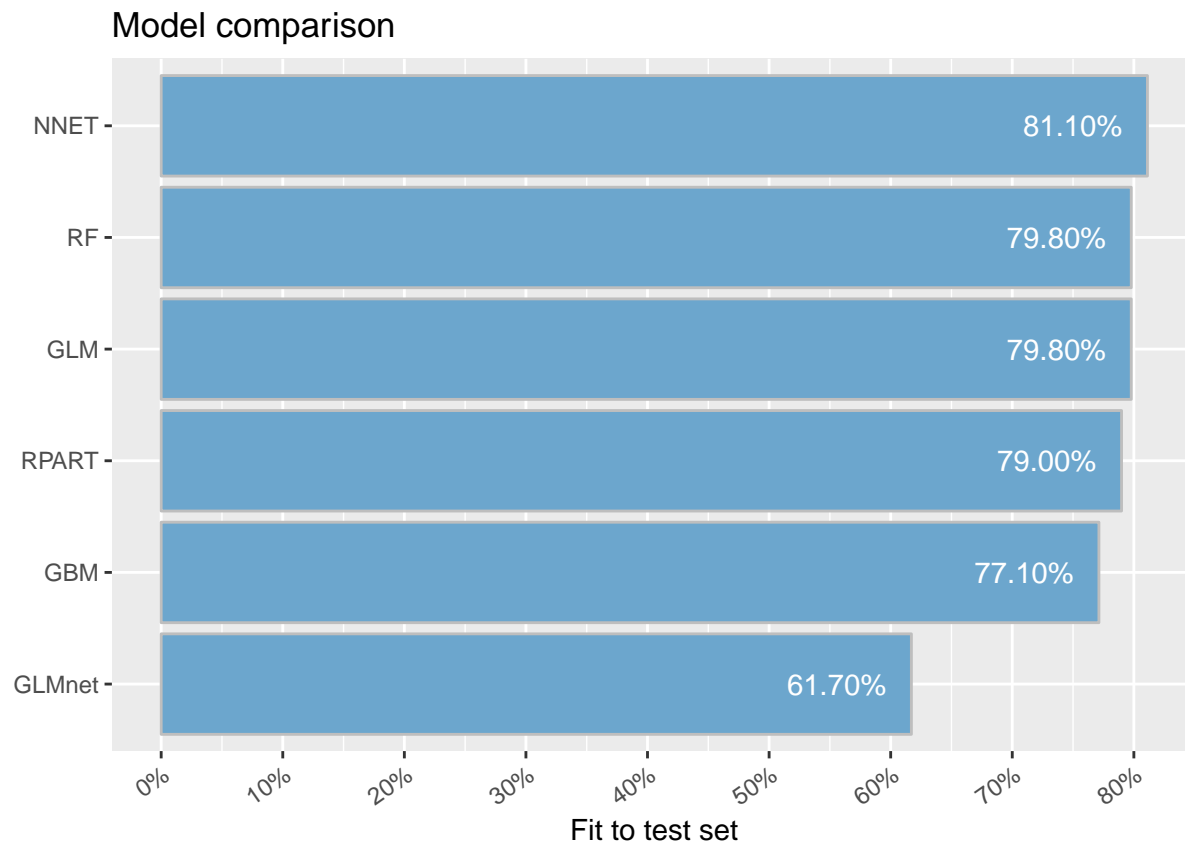
After the RFE, we retain 12 predictors: Country, SS, Ethnicity, Oscore, Age, Cscore, Education, Gender, Impulsive, Escore, Ascore, Nscore



The comparative analysis of the contribution of each factor agrees by and large with that of the density distribution plots: among the personality trait tests, N-score and E-score contribute the least while seeking sensation and O-score contribute the most. However, the impulsivity doesn't seem to be as strong a factor as one might have expected from the density plots.

Model comparisons





Results

GBM: Little importance is given to the education variable. Improvement analysis -> what did the use of ML contribute

Demographics have more impact than personality

Conclusion

Scope -> having used once in lifetime may not be telling. Could be interesting to bin never with over ten years ago ?